

PROJECT : Insurance Claims- Fraud Detection

```
In [36]: #Importing the libreris like pandas, numpy for seleceng the data and converng the data to
import pandas as pd
import numpy as np
import warnings
warnings.filterwarnings('ignore')
```

```
In [37]: #Creating the variable df and loading the dataset to the variable df
df=pd.read_csv('data')
df
```

Out[37]:

	months_as_customer	age	policy_number	policy_bind_date	policy_state	policy_csl	policy_deductable	poli
0	328	48	521585	17-10-2014	OH	250/500	1000	
1	228	42	342868	27-06-2006	IN	250/500	2000	
2	134	29	687698	06-09-2000	OH	100/300	2000	
3	256	41	227811	25-05-1990	IL	250/500	2000	
4	228	44	367455	06-06-2014	IL	500/1000	1000	
...
995	3	38	941851	16-07-1991	OH	500/1000	1000	
996	285	41	186934	05-01-2014	IL	100/300	1000	
997	130	34	918516	17-02-2003	OH	250/500	500	
998	458	62	533940	18-11-2011	IL	500/1000	2000	
999	456	60	556080	11-11-1996	OH	250/500	1000	

1000 rows × 40 columns

```
In [4]: # let's get the information about the dataset

df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1000 entries, 0 to 999
Data columns (total 40 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   months_as_customer                    1000 non-null   int64
1   age                                   1000 non-null   int64
2   policy_number                        1000 non-null   int64
3   policy_bind_date                     1000 non-null   object
4   policy_state                         1000 non-null   object
5   policy_csl                           1000 non-null   object
6   policy_deductable                    1000 non-null   int64
7   policy_annual_premium                1000 non-null   float64
8   umbrella_limit                       1000 non-null   int64
9   insured_zip                          1000 non-null   int64
10  insured_sex                          1000 non-null   object
11  insured_education_level              1000 non-null   object
12  insured_occupation                   1000 non-null   object
13  insured_hobbies                      1000 non-null   object
Loading [MathJax]/extensions/Safe.js relationship 1000 non-null   object
```

15	capital-gains	1000	non-null	int64
16	capital-loss	1000	non-null	int64
17	incident_date	1000	non-null	object
18	incident_type	1000	non-null	object
19	collision_type	1000	non-null	object
20	incident_severity	1000	non-null	object
21	authorities_contacted	1000	non-null	object
22	incident_state	1000	non-null	object
23	incident_city	1000	non-null	object
24	incident_location	1000	non-null	object
25	incident_hour_of_the_day	1000	non-null	int64
26	number_of_vehicles_involved	1000	non-null	int64
27	property_damage	1000	non-null	object
28	bodily_injuries	1000	non-null	int64
29	witnesses	1000	non-null	int64
30	police_report_available	1000	non-null	object
31	total_claim_amount	1000	non-null	int64
32	injury_claim	1000	non-null	int64
33	property_claim	1000	non-null	int64
34	vehicle_claim	1000	non-null	int64
35	auto_make	1000	non-null	object
36	auto_model	1000	non-null	object
37	auto_year	1000	non-null	int64
38	fraud_reported	1000	non-null	object
39	_c39	0	non-null	float64

dtypes: float64(2), int64(17), object(21)
memory usage: 312.6+ KB

```
In [38]: # let's check whether the data has any null values or not.

# but there is '?' in the dataset which we have to replace by NaN Values
df = df.replace('?', np.NaN)

df.isnull().any()
```

```
Out[38]: months_as_customer      False
age                             False
policy_number                   False
policy_bind_date                False
policy_state                    False
policy_csl                      False
policy_deductable               False
policy_annual_premium           False
umbrella_limit                  False
insured_zip                     False
insured_sex                     False
insured_education_level         False
insured_occupation              False
insured_hobbies                 False
insured_relationship            False
capital-gains                   False
capital-loss                    False
incident_date                   False
incident_type                   False
collision_type                  True
incident_severity               False
authorities_contacted           False
incident_state                  False
incident_city                   False
incident_location               False
incident_hour_of_the_day        False
number_of_vehicles_involved     False
property_damage                 True
Loading [MathJax]/extensions/Safe.js S False
```

witnesses	False
police_report_available	True
total_claim_amount	False
injury_claim	False
property_claim	False
vehicle_claim	False
auto_make	False
auto_model	False
auto_year	False
fraud_reported	False
_c39	True

dtype: bool

```
In [39]: # filling the null values

# we will replace the '?' by the most common collision type as we are unaware of the type.
df['collision_type'].fillna(df['collision_type'].mode()[0], inplace = True)

# It may be the case that there are no responses for property damage then we might take it
df['property_damage'].fillna('NO', inplace = True)

# again, if there are no responses fpr police report available then we might take it as No
df['police_report_available'].fillna('NO', inplace = True)
```

```
In [40]: df.isnull().sum()
```

```
Out[40]: months_as_customer    0
age                            0
policy_number                 0
policy_bind_date              0
policy_state                  0
policy_csl                    0
policy_deductable             0
policy_annual_premium         0
umbrella_limit                0
insured_zip                   0
insured_sex                   0
insured_education_level       0
insured_occupation            0
insured_hobbies                0
insured_relationship           0
capital-gains                 0
capital-loss                  0
incident_date                 0
incident_type                 0
collision_type                0
incident_severity             0
authorities_contacted         0
incident_state                0
incident_city                 0
incident_location             0
incident_hour_of_the_day      0
number_of_vehicles_involved   0
property_damage               0
bodily_injuries               0
witnesses                     0
police_report_available       0
total_claim_amount            0
injury_claim                  0
property_claim                0
vehicle_claim                 0
auto_make                     0
```

```
auto_year      0
fraud_reported 0
_c39           1000
dtype: int64
```

```
In [41]: #dropping the _c39 column, _c39 column doesnot have any impact on target column
df=df.drop(columns='_c39')
df
```


```
Out[41]:
```

	months_as_customer	age	policy_number	policy_bind_date	policy_state	policy_csl	policy_deductable	poli
0	328	48	521585	17-10-2014	OH	250/500	1000	
1	228	42	342868	27-06-2006	IN	250/500	2000	
2	134	29	687698	06-09-2000	OH	100/300	2000	
3	256	41	227811	25-05-1990	IL	250/500	2000	
4	228	44	367455	06-06-2014	IL	500/1000	1000	
...
995	3	38	941851	16-07-1991	OH	500/1000	1000	
996	285	41	186934	05-01-2014	IL	100/300	1000	
997	130	34	918516	17-02-2003	OH	250/500	500	
998	458	62	533940	18-11-2011	IL	500/1000	2000	
999	456	60	556080	11-11-1996	OH	250/500	1000	

1000 rows × 39 columns

```
In [9]: import matplotlib.pyplot as plt
import seaborn as sns
plt.style.use('fivethirtyeight')
ax = sns.countplot(x='fraud_reported', data=df, hue='fraud_reported')
df['fraud_reported'].value_counts()
```


```
Out[9]: N    753
Y     247
Name: fraud_reported, dtype: int64
```



```
In [10]: print(df['incident_state'].value_counts())
plt.style.use('fivethirtyeight')
figure(figsize=(10,6))
```


```
ax = df.groupby('incident_state').fraud_reported.count().plot.bar(ylim=0)
ax.set_ylabel('Fraud reported')
plt.show()
```

```
NY    262
SC    248
WV    217
VA    110
NC    110
PA     30
OH     23
Name: incident_state, dtype: int64
```




In [11]:


```
plt.style.use('fivethirtyeight')
fig = plt.figure(figsize=(18,6))
ax = df.groupby('incident_date').total_claim_amount.count().plot.bar(ylim=0)
ax.set_ylabel('Claim amount ($)')
plt.show()
```



```
In [12]: plt.style.use('fivethirtyeight')
fig = plt.figure(figsize=(10,6))
ax = df.groupby('incident_type').fraud_reported.count().plot.bar(ylim=0)
ax.set_xticklabels(ax.get_xticklabels(), rotation=20, ha="right")
ax.set_ylabel('Fraud reported')
plt.show()
```




```
In [13]: fig = plt.figure(figsize=(10,6))
ax = sns.countplot(y = 'insured_education_level', data=df)
ax.set_ylabel('policy_annual_premium')
plt.show()
```




In [14]:

```
fig = plt.figure(figsize=(10,6))
ax = (df['insured_sex'].value_counts()*100.0 /len(df))\
.plot.pie(autopct='%1f%%', labels = ['Male', 'Female'], fontsize=12)
ax.set_title('% Gender')
plt.show()
```




In [15]:


```
fig = plt.figure(figsize=(10,6))
ax = (df['insured_relationship'].value_counts()*100.0 /len(df))\
.plot.pie(autopct='%1f%%', labels = ['husband', 'wife', 'own-child', 'unmarried', 'other-
      ], fontsize=12)
ax.set_title('% Relationship')
plt.show()
```



```
In [16]: fig = plt.figure(figsize=(10,6))
ax = (df['authorities_contacted'].value_counts()*100.0 /len(df))\
.plot.pie(autopct='%.1f%%', labels = ['Police', 'Fire', 'Other', 'None', 'Ambulance'],
         fontsize=12)
```



```
In [17]: fig = plt.figure(figsize=(10,6))
ax = (df['incident_severity'].value_counts()*100.0 /len(df))\
.plot.pie(autopct='%.1f%%', labels = ['Major Damage', 'Total Loss', 'Minor Damage', 'Trivial'],
         fontsize=12)
```

```
In [19]: # let's extract days, month and year from policy bind date

df['policy_bind_date'] = pd.to_datetime(df['policy_bind_date'], errors = 'coerce')
df
```

```
Out[19]:
```

	months_as_customer	age	policy_number	policy_bind_date	policy_state	policy_csl	policy_deductable	poli
0	328	48	521585	2014-10-17	OH	250/500	1000	
1	228	42	342868	2006-06-27	IN	250/500	2000	
2	134	29	687698	2000-06-09	OH	100/300	2000	
3	256	41	227811	1990-05-25	IL	250/500	2000	
4	228	44	367455	2014-06-06	IL	500/1000	1000	
...
995	3	38	941851	1991-07-16	OH	500/1000	1000	
996	285	41	186934	2014-05-01	IL	100/300	1000	
997	130	34	918516	2003-02-17	OH	250/500	500	
998	458	62	533940	2011-11-18	IL	500/1000	2000	
999	456	60	556080	1996-11-11	OH	250/500	1000	

1000 rows × 39 columns

```
In [42]: # let's encode the fraud report to numerical values

df['fraud_reported'] = df['fraud_reported'].replace(('Y', 'N'), (0, 1))
```

```
In [25]: df.columns
```

```
Out[25]: Index(['months_as_customer', 'age', 'policy_number', 'policy_bind_date',
      'policy_state', 'policy_csl', 'policy_deductable',
      'policy_annual_premium', 'umbrella_limit', 'insured_zip', 'insured_sex',
      'insured_education_level', 'insured_occupation', 'insured_hobbies',
```

```
'insured_relationship', 'capital-gains', 'capital-loss',
'incident_date', 'incident_type', 'collision_type', 'incident_severity',
'authorities_contacted', 'incident_state', 'incident_city',
'incident_location', 'incident_hour_of_the_day',
'number_of_vehicles_involved', 'property_damage', 'bodily_injuries',
'witnesses', 'police_report_available', 'total_claim_amount',
'injury_claim', 'property_claim', 'vehicle_claim', 'auto_make',
'auto_model', 'auto_year', 'fraud_reported'],
dtype='object')
```

In [43]:

```
# let's check the correlation of authorities_contacted with the target
# changing all the categorical column into numerical columns

df[['auto_model', 'fraud_reported']].groupby(['auto_model'],
                                              as_index = False).mean().sort_values(by = 'fraud_reported', ascending = False)
```

Out[43]:

	auto_model	fraud_reported
0	3 Series	0.944444
31	RSX	0.916667
25	Malibu	0.900000
36	Wrangler	0.880952
29	Pathfinder	0.870968
35	Ultima	0.869565
9	Camry	0.857143
11	Corolla	0.850000
8	CRV	0.850000
21	Legacy	0.843750
27	Neon	0.837838
3	95	0.814815
33	TL	0.800000
2	93	0.800000
23	MDX	0.777778
6	Accord	0.769231
17	Grand Cherokee	0.760000
13	Escape	0.750000
12	E400	0.740741
4	A3	0.729730
18	Highlander	0.727273
28	Passat	0.727273
1	92x	0.714286
20	Jetta	0.714286
16	Fusion	0.714286
15	Forrestor	0.714286
26	Maxima	0.708333
19	Impreza	0.700000

5 0.695652

	auto_model	fraud_reported
30	RAM	0.674419
22	M5	0.666667
5	A5	0.656250
10	Civic	0.636364
14	F150	0.629630
34	Tahoe	0.625000
7	C300	0.611111
24	ML350	0.600000
32	Silverado	0.590909
38	X6	0.562500

In [113]:

```
# let's perform target encoding for auto make

df['auto_model'] = df['auto_model'].replace(('3 Series', 'RSX', 'Malibu', 'Wrangler', 'Pathfinder', 'Corolla', 'CRV', 'Legacy', 'Neon', '95', 'TL', '93', 'MDX', 'Accord', 'Grand Cherokee', 'A3', 'Highlander', 'Passat', '92x', 'Jetta', 'Fusion', 'Forrester', 'Maxima', 'Impreza', 'Civic', 'F150', 'Tahoe', 'C300', 'ML350', 'Silverado', 'X6'),
      (0.95, 0.91, 0.90, 0.88, 0.87, 0.86, 0.855, 0.85, 0.85, 0.84, 0.83, 0.81, 0.80, 0.80, 0.79, 0.78, 0.77, 0.76, 0.75, 0.74, 0.73, 0.72, 0.72, 0.71, 0.71, 0.71, 0.71, 0.70, 0.70, 0.69, 0.67, 0.66, 0.65, 0.64, 0.63, 0.62, 0.61, 0.60, 0.59, 0.58, 0.57, 0.56, 0.55, 0.54, 0.53, 0.52, 0.51, 0.50, 0.49, 0.48, 0.47, 0.46, 0.45, 0.44, 0.43, 0.42, 0.41, 0.40, 0.39, 0.38, 0.37, 0.36, 0.35, 0.34, 0.33, 0.32, 0.31, 0.30, 0.29, 0.28, 0.27, 0.26, 0.25, 0.24, 0.23, 0.22, 0.21, 0.20, 0.19, 0.18, 0.17, 0.16, 0.15, 0.14, 0.13, 0.12, 0.11, 0.10, 0.09, 0.08, 0.07, 0.06, 0.05, 0.04, 0.03, 0.02, 0.01, 0.00))
```

```
-----
KeyError                                Traceback (most recent call last)
C:\ProgramData\Anaconda3\lib\site-packages\pandas\core\indexes\base.py in get_loc(self, key, method, tolerance)
    3360         try:
-> 3361             return self._engine.get_loc(casted_key)
    3362         except KeyError as err:

C:\ProgramData\Anaconda3\lib\site-packages\pandas\_libs\index.pyx in pandas._libs.index.IndexEngine.get_loc()

C:\ProgramData\Anaconda3\lib\site-packages\pandas\_libs\index.pyx in pandas._libs.index.IndexEngine.get_loc()

pandas\_libs\hashtable_class_helper.pxi in pandas._libs.hashtable.PyObjectHashTable.get_item()

pandas\_libs\hashtable_class_helper.pxi in pandas._libs.hashtable.PyObjectHashTable.get_item()

KeyError: 'auto_model'
```

The above exception was the direct cause of the following exception:

```
KeyError                                Traceback (most recent call last)
~\AppData\Local\Temp\ipykernel_8104\2100930274.py in <module>
      1 # let's perform target encoding for auto make
      2
----> 3 df['auto_model'] = df['auto_model'].replace(('3 Series', 'RSX', 'Malibu', 'Wrangler', 'Pathfinder', 'Ultima', 'Camry', 'Corolla', 'CRV', 'Legacy', 'Neon', '95', 'TL', '93', 'MDX', 'Accord', 'Grand Cherokee', 'Escape', 'E400', 'A3', 'Highlander', 'Passat', '92x', 'Jetta', 'Fusion', 'Forrester', 'Maxima', 'Impreza', 'X5', 'RAM', 'M5', 'A5', 'Civic', 'F150', 'Tahoe', 'C300', 'ML350', 'Silverado', 'X6'),
      4         (0.95, 0.91, 0.90, 0.88, 0.87, 0.86, 0.855, 0.85, 0.85, 0.84, 0.83, 0.81, 0.80, 0.80, 0.79, 0.78, 0.77, 0.76, 0.75, 0.74, 0.73, 0.72, 0.72, 0.71, 0.71, 0.71, 0.71, 0.70, 0.70, 0.69, 0.67, 0.66, 0.65, 0.64, 0.63, 0.62, 0.61, 0.60, 0.59, 0.58, 0.57, 0.56, 0.55, 0.54, 0.53, 0.52, 0.51, 0.50, 0.49, 0.48, 0.47, 0.46, 0.45, 0.44, 0.43, 0.42, 0.41, 0.40, 0.39, 0.38, 0.37, 0.36, 0.35, 0.34, 0.33, 0.32, 0.31, 0.30, 0.29, 0.28, 0.27, 0.26, 0.25, 0.24, 0.23, 0.22, 0.21, 0.20, 0.19, 0.18, 0.17, 0.16, 0.15, 0.14, 0.13, 0.12, 0.11, 0.10, 0.09, 0.08, 0.07, 0.06, 0.05, 0.04, 0.03, 0.02, 0.01, 0.00))
      5
```

```

C:\ProgramData\Anaconda3\lib\site-packages\pandas\core\frame.py in __getitem__(self, key)
    3456         if self.columns.nlevels > 1:
    3457             return self._getitem_multilevel(key)
-> 3458         indexer = self.columns.get_loc(key)
    3459         if is_integer(indexer):
    3460             indexer = [indexer]

C:\ProgramData\Anaconda3\lib\site-packages\pandas\core\indexes\base.py in get_loc(self, key, method, tolerance)
    3361         return self._engine.get_loc(casted_key)
    3362     except KeyError as err:
-> 3363         raise KeyError(key) from err
    3364
    3365     if is_scalar(key) and isna(key) and not self.hasnans:

KeyError: 'auto_model'

```

```

In [51]: #let's check the correlation auto make with the target

df[['auto_make', 'fraud_reported']].groupby(['auto_make'],
                                             as_index = False).mean().sort_values(by = 'fraud_reported', ascending = False)

```

```

Out[51]:
   auto_make  fraud_reported
12         0.84         0.835821
11         0.82         0.820513
10         0.81         0.814286
9          0.80         0.808824
8          0.77         0.775000
7          0.76         0.762500
6          0.75         0.750000
5          0.74         0.745455
4          0.73         0.723684
3          0.72         0.722222
2          0.71         0.720588
1          0.69         0.695035
0          0.66         0.661538

```

```

In [52]: # let's perform target encoding for auto make

df['auto_make'] = df['auto_make'].replace(('Jeep', 'Nissan', 'Toyota', 'Accura', 'Saab', 'Subur',
                                           'Dodge', 'Honda', 'Chevrolet', 'BMW', 'Volkswagen', 'Audi', 'Ford',
                                           0.84, 0.82, 0.81, 0.80, 0.77, 0.76, 0.75, 0.74, 0.73, 0.72, 0.71, 0.69, 0.66))

```

```

In [54]: # changing all the categorical column into numerical columns

df[['police_report_available', 'fraud_reported']].groupby(['police_report_available'],
                                                           as_index = False).mean().sort_values(by = 'fraud_reported', ascending = False)

```

```

Out[54]:
   police_report_available  fraud_reported
1                YES         0.770701
0                NO         0.744898

```

```
In [56]: df['police_report_available'] = df['police_report_available'].replace(('NO', 'YES'), (0.77, 0.74))
```

```
In [57]: df[['property_damage', 'fraud_reported']].groupby(['property_damage'],
    as_index = False).mean().sort_values(by = 'fraud_reported', ascending = False)
```

Out[57]:

	property_damage	fraud_reported
0	NO	0.757880
1	YES	0.741722

```
In [58]: df['property_damage'] = df['property_damage'].replace(('NO', 'YES'), (0.76, 0.74))
```

```
In [60]: df[['incident_city', 'fraud_reported']].groupby(['incident_city'],
    as_index = False).mean().sort_values(by = 'fraud_reported', ascending = False)
```

Out[60]:

	incident_city	fraud_reported
4	Northbrook	0.778689
5	Riverwood	0.776119
3	Northbend	0.765517
6	Springfield	0.757962
2	Hillsdale	0.751773
1	Columbus	0.738255
0	Arlington	0.710526

```
In [63]: df['incident_city'] = df['incident_city'].replace(('Northbrook', 'Riverwood', 'Northbend', 'Springfield', 'Hillsdale', 'Columbus', 'Arlington'), (0.78, 0.77, 0.76, 0.75, 0.74, 0.73, 0.72))
```

```
In [64]: df[['incident_state', 'fraud_reported']].groupby(['incident_state'],
    as_index = False).mean().sort_values(by = 'fraud_reported', ascending = False)
```

Out[64]:

	incident_state	fraud_reported
6	WV	0.820276
1	NY	0.778626
5	VA	0.772727
3	PA	0.733333
4	SC	0.705645
0	NC	0.690909
2	OH	0.565217

```
In [65]: df['incident_state'] = df['incident_state'].replace(('WV', 'NY', 'VA', 'PA', 'SC', 'NC', 'OH'), (0.82, 0.77, 0.76, 0.73, 0.70, 0.69, 0.57))
```

```
df[['authorities_contacted', 'fraud_reported']].groupby(['authorities_contacted'],
as_index = False).mean().sort_values(by = 'fraud_reported', ascending = False)
```

Out[66]:

	authorities_contacted	fraud_reported
2	None	0.934066
4	Police	0.791096
1	Fire	0.730942
0	Ambulance	0.709184
3	Other	0.681818

In [69]:

```
df['authorities_contacted'] = df['authorities_contacted'].replace(('None', 'Police', 'Fire', 'Ambulance', 'Other'),
(0.94,0.79,0.73,0.70,0.68))
```

In [70]:

```
df[['incident_severity', 'fraud_reported']].groupby(['incident_severity'],
as_index = False).mean().sort_values(by = 'fraud_reported', ascending = False)
```

Out[70]:

	incident_severity	fraud_reported
3	Trivial Damage	0.933333
1	Minor Damage	0.892655
2	Total Loss	0.871429
0	Major Damage	0.394928

In [71]:

```
df['incident_severity'] = df['incident_severity'].replace(('Trivial Damage', 'Minor Damage', 'Total Loss', 'Major Damage'),
(0.94,0.89,0.87,0.39))
```

In [72]:

```
df[['collision_type', 'fraud_reported']].groupby(['collision_type'],
as_index = False).mean().sort_values(by = 'fraud_reported', ascending = False)
```

Out[72]:

	collision_type	fraud_reported
1	Rear Collision	0.772340
2	Side Collision	0.746377
0	Front Collision	0.724409

In [73]:

```
df['collision_type'] = df['collision_type'].replace(('Rear Collision', 'Side Collision', 'Front Collision'),
(0.78,0.74,0.72))
```

In [74]:

```
df[['incident_type', 'fraud_reported']].groupby(['incident_type'],
as_index = False).mean().sort_values(by = 'fraud_reported', ascending = False)
```

Out[74]:

	incident_type	fraud_reported
3	Vehicle Theft	0.914894
1	Parked Car	0.904762
0	Multi-vehicle Collision	0.727924
2	Single-vehicle Collision	0.709677

```
In [75]: df['incident_type'] = df['incident_type'].replace(('Vehicle Theft','Parked Car','Multi-veh  
          'Single Vehicle Collision')),(0.91, 0.90, 0.72,0.70))
```

```
In [76]: df['incident_date'] = pd.to_datetime(df['incident_date'], errors = 'coerce')  
  
# extracting days and month from date  
df['incident_month'] = df['incident_date'].dt.month  
df['incident_day'] = df['incident_date'].dt.day
```

```
In [77]: df[['insured_relationship','fraud_reported']].groupby(['insured_relationship'],  
          as_index = False).mean().sort_values(by = 'fraud_reported', ascending = Fa
```

Out[77]:

	insured_relationship	fraud_reported
0	husband	0.794118
3	own-child	0.786885
4	unmarried	0.758865
1	not-in-family	0.741379
5	wife	0.729032
2	other-relative	0.706215

```
In [78]: df['insured_relationship'] = df['insured_relationship'].replace(('husband','own-child','ur  
          'not-in-family','wife','other-relative')),(0.79,0.7
```

```
In [79]: df[['insured_hobbies','fraud_reported']].groupby(['insured_hobbies'],  
          as_index = False).mean().sort_values(by = 'fraud_reported', ascending = Fa
```

Out[79]:

	insured_hobbies	fraud_reported
4	camping	0.909091
11	kayaking	0.907407
9	golf	0.890909
7	dancing	0.883721
3	bungie-jumping	0.839286
12	movies	0.836364
1	basketball	0.823529
8	exercise	0.807018
17	sleeping	0.804878
18	video-games	0.800000
16	skydiving	0.775510
13	paintball	0.771930
10	hiking	0.769231
0	base-jumping	0.734694
15	reading	0.734375
	Loading [MathJax]/extensions/Safe.js olo	0.723404

	insured_hobbies	fraud_reported
2	board-games	0.708333
19	yachting	0.698113
6	cross-fit	0.257143
5	chess	0.173913

In [80]:

```
df['insured_hobbies'] = df['insured_hobbies'].replace(('camping', 'kayaking', 'golf', 'danc
    'bungee-jumping', 'movies', 'basketball', 'exercise', 'sleeping', 'video-games', 'skydi
    'hiking', 'base-jumping', 'reading', 'polo', 'board-games', 'yachting', 'cross-fit'
    0.89, 0.88, 0.84, 0.83, 0.82, 0.81, 0.805, 0.80, 0.78, 0.77, 0.76, 0.73, 0.73, 0.72, 0.
```

In [81]:

```
df[['insured_occupation', 'fraud_reported']].groupby(['insured_occupation'],
    as_index = False).mean().sort_values(by = 'fraud_reported', ascending = Fa
```

Out[81]:

	insured_occupation	fraud_reported
7	other-service	0.830986
8	priv-house-serv	0.830986
0	adm-clerical	0.830769
5	handlers-cleaners	0.796296
9	prof-specialty	0.788235
10	protective-serv	0.777778
6	machine-op-inspct	0.763441
1	armed-forces	0.753623
11	sales	0.723684
12	tech-support	0.717949
13	transport-moving	0.708333
2	craft-repair	0.702703
4	farming-fishing	0.698113
3	exec-managerial	0.631579

In [82]:

```
df['insured_occupation'] = df['insured_occupation'].replace(('other-service', 'priv-house-s
    'adm-clerical', 'handlers-cleaners', 'prof-specialty', 'protective-se
    'machine-op-inspct', 'armed-forces', 'sales', 'tech-support', 'transport-movin
    'farming-fishing', 'exec-managerial'), (0.84, 0.84, 0.83, 0.79, 0.78, 0.77,
    0.705, 0.70, 0.69, 0.63))
```

In [83]:

```
df[['insured_education_level', 'fraud_reported']].groupby(['insured_education_level'],
    as_index = False).mean().sort_values(by = 'fraud_reported', ascending = Fa
```

Out[83]:

	insured_education_level	fraud_reported
5	Masters	0.776224
2	High School	0.775000
0	Associate	0.765517
	JD	0.739130

	insured_education_level	fraud_reported
1	College	0.737705
4	MD	0.736111
6	PhD	0.736000

```
In [84]: df['insured_education_level'] = df['insured_education_level'].replace(('Masters', 'High School', 'JD', 'College', 'MD', 'PhD'), (0.78, 0.77, 0.76, 0.74, 0.73, 0.72))
```

```
In [85]: df[['insured_sex', 'fraud_reported']].groupby(['insured_sex'], as_index = False).mean().sort_index(by = 'fraud_reported', ascending = False)
```

```
Out[85]:
```

	insured_sex	fraud_reported
0	FEMALE	0.765363
1	MALE	0.738661

```
In [86]: df['insured_sex'] = df['insured_sex'].replace(('FEMALE', 'MALE'), (0.76, 0.73))
```

```
In [87]: df[['policy_csl', 'fraud_reported']].groupby(['policy_csl'], as_index = False).mean().sort_index(by = 'fraud_reported', ascending = False)
```

```
Out[87]:
```

	policy_csl	fraud_reported
2	500/1000	0.783333
0	100/300	0.742120
1	250/500	0.737892

```
In [88]: df['policy_csl'] = df['policy_csl'].replace(('500/1000', '100/300', '250/500'), (0.78, 0.74, 0.73))
```

```
In [89]: df[['policy_state', 'fraud_reported']].groupby(['policy_state'], as_index = False).mean().sort_index(by = 'fraud_reported', ascending = False)
```

```
Out[89]:
```

	policy_state	fraud_reported
0	IL	0.772189
1	IN	0.745161
2	OH	0.741477

```
In [90]: df['policy_state'] = df['policy_state'].replace(('IL', 'IN', 'OH'), (0.77, 0.745, 0.74))
```

```
In [91]: # let's delete unnecassary columns

df = df.drop(['policy_number', 'policy_bind_date', 'incident_date', 'incident_location', 'auto_policy'])

# let's check the columns after deleting the columns
df.columns
```

```
Out[91]: Index(['months_as_customer', 'age', 'policy_state', 'policy_csl',
              'policy_deductable', 'policy_annual_premium', 'umbrella_limit',
              'insured_zip', 'insured_sex', 'insured_education_level',
              'insured_occupation', 'insured_hobbies', 'insured_relationship',
              'capital-gains', 'capital-loss', 'incident_type', 'collision_type',
              'incident_severity', 'authorities_contacted', 'incident_state',
              'incident_city', 'incident_hour_of_the_day',
              'number_of_vehicles_involved', 'property_damage', 'bodily_injuries',
              'witnesses', 'police_report_available', 'total_claim_amount',
              'injury_claim', 'property_claim', 'vehicle_claim', 'auto_make',
              'auto_year', 'fraud_reported', 'incident_month', 'incident_day'],
              dtype='object')
```


```
In [127... # now all the categorical columns are converted into numerical columns
# All the unnecassary columns also dropped from df dataset
# now lets see the df dataset
df
```

```
Out[127... months_as_customer  age  policy_state  policy_csl  policy_deductable  policy_annual_premium  umbrella_limit
```

0	328	48	0.740	0.73	1000	1406.91	0
1	228	42	0.745	0.73	2000	1197.22	5000000
2	134	29	0.740	0.74	2000	1413.14	5000000
3	256	41	0.770	0.73	2000	1415.74	6000000
4	228	44	0.770	0.78	1000	1583.91	6000000
...
995	3	38	0.740	0.78	1000	1310.80	0
996	285	41	0.770	0.74	1000	1436.79	0
997	130	34	0.740	0.73	500	1383.49	3000000
998	458	62	0.770	0.78	2000	1356.92	5000000
999	456	60	0.740	0.73	1000	766.19	0

1000 rows × 36 columns

```
In [129... # checking the co-relation values with the help of heat map
import matplotlib.pyplot as plt
corrmat = df.corr()
plt.figure(figsize=(30,30))
#plot heat map
g=sns.heatmap(corrmat,annot=True)
```



In [92]:

```
# let's split the data into dependent and independent sets

x = df.drop(['fraud_reported'], axis = 1)
y = df['fraud_reported']
# checking the x, y shape
print("Shape of x :", x.shape)
print("Shape of y :", y.shape)
```

Shape of x : (1000, 35)
Shape of y : (1000,)

In [93]:

```
#removing the skewness by yeo johnson method
from sklearn.preprocessing import power_transform
x=power_transform(x,method='yeo-johnson')
x
```

array([[1.05127872, 1.00873272, 0. , ..., -0.19745541, 0.045698, 1.04604418],

```
[ 0.30453584, 0.43143333, 0.      , ..., 0.30262249,
 -1.10045698, 0.82475977],
 [-0.51122603, -1.13951302, 0.      , ..., 0.30262249,
 -0.02541209, 0.88243021],
 ...,
 [-0.54970642, -0.47044379, 0.      , ..., -1.49627768,
 -1.10045698, 0.93846168],
 [ 1.9126537 , 2.13055434, 0.      , ..., -1.17624558,
 -0.02541209, 1.09778562],
 [ 1.90009963, 1.98572613, 0.      , ..., 0.30262249,
 -0.02541209, 1.09778562]])
```

```
In [94]: pd.DataFrame(x).skew()
```

```
Out[94]: 0    -0.135661
1    -0.001945
2     0.000000
3     0.000000
4     0.023988
5     0.004758
6    -7.865930
7     0.000000
8    -0.148630
9    -0.002118
10   -0.020558
11   -0.294249
12   -0.020572
13    0.038722
14    0.090488
15    0.941516
16   -0.072601
17   -0.828981
18    0.231201
19   -0.038724
20   -0.049267
21   -0.256957
22    0.363693
23   -0.863806
24   -0.128799
25   -0.153648
26   -0.802728
27   -0.510354
28   -0.415781
29   -0.358814
30   -0.522718
31    0.009484
32   -0.012491
33    0.306468
34   -0.221642
dtype: float64
```

```
In [95]: from sklearn.preprocessing import MinMaxScaler
mms=MinMaxScaler()
x1=mms.fit_transform(x)
```

```
In [96]: # importing all the algorithms for checking the accuracy_score and model performance
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import train_test_split,cross_val_score,GridSearchCV
from sklearn.metrics import accuracy_score,confusion_matrix,classification_report,f1_score
from sklearn.tree import DecisionTreeClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.svm import SVC
```

```
from sklearn.naive_bayes import GaussianNB
from sklearn.ensemble import RandomForestClassifier
```

In [98]:

```
model=[RandomForestClassifier(),DecisionTreeClassifier(),SVC(),KNeighborsClassifier(),GaussianNB()]
max_r2_score=0
for i_state in range(0,10):
    x_train,x_test,y_train,y_test=train_test_split(x1,y,random_state=i_state,test_size=0.2)
    for a in model:
        a.fit(x_train,y_train)
        pred=a.predict(x_test)
        score=accuracy_score(y_test,pred)
        print('score for random_state',i_state,'is',score)
        if score>max_r2_score:
            max_r2_score=score
            Final_state=i_state
            Final_model= a
print('accuracy_score ',max_r2_score,'for random state ',Final_state, 'and model is ',Final_model)
```

```
score for random_state 0 is 0.785
score for random_state 0 is 0.83
score for random_state 0 is 0.81
score for random_state 0 is 0.7
score for random_state 0 is 0.795
score for random_state 0 is 0.84
score for random_state 1 is 0.845
score for random_state 1 is 0.775
score for random_state 1 is 0.82
score for random_state 1 is 0.75
score for random_state 1 is 0.84
score for random_state 1 is 0.84
score for random_state 2 is 0.83
score for random_state 2 is 0.855
score for random_state 2 is 0.855
score for random_state 2 is 0.76
score for random_state 2 is 0.865
score for random_state 2 is 0.87
score for random_state 3 is 0.845
score for random_state 3 is 0.745
score for random_state 3 is 0.825
score for random_state 3 is 0.75
score for random_state 3 is 0.83
score for random_state 3 is 0.865
score for random_state 4 is 0.805
score for random_state 4 is 0.775
score for random_state 4 is 0.805
score for random_state 4 is 0.685
score for random_state 4 is 0.84
score for random_state 4 is 0.83
score for random_state 5 is 0.835
score for random_state 5 is 0.78
score for random_state 5 is 0.81
score for random_state 5 is 0.755
score for random_state 5 is 0.805
score for random_state 5 is 0.815
score for random_state 6 is 0.85
score for random_state 6 is 0.79
score for random_state 6 is 0.845
score for random_state 6 is 0.795
score for random_state 6 is 0.795
score for random_state 7 is 0.83
score for random_state 7 is 0.76
score for random_state 7 is 0.79
```



```

score for random_state 7 is 0.72
score for random_state 7 is 0.77
score for random_state 7 is 0.815
score for random_state 8 is 0.83
score for random_state 8 is 0.73
score for random_state 8 is 0.835
score for random_state 8 is 0.785
score for random_state 8 is 0.815
score for random_state 8 is 0.855
score for random_state 9 is 0.84
score for random_state 9 is 0.75
score for random_state 9 is 0.81
score for random_state 9 is 0.74
score for random_state 9 is 0.77
score for random_state 9 is 0.835
accuracy_score 0.87 for random state 2 and model is LogisticRegression()

```

In [99]:

```

# we are training the model with LogisticRegression for randomstate 2 and checking the acc
lr=LogisticRegression()
x_train,x_test,y_train,y_test=train_test_split(x1,y,random_state=2,test_size=0.2)
lr.fit(x_train,y_train)
lr.score(x_train,y_train)
pred_y=lr.predict(x_test)
lrs=accuracy_score(y_test,pred_y)
print('accuracy_score =',lrs*100)
print(classification_report(y_test,pred_y))
print(confusion_matrix(y_test,pred_y))
print('F1_score = ',f1_score(y_test,pred_y)*100)
from sklearn.model_selection import cross_val_score
cv_score=cross_val_score(lr,x1,y,cv=5)
cv_mean=cv_score.mean()
print("cross_val_score=",cv_mean*100)

```

```

accuracy_score = 87.0
          precision    recall  f1-score   support

      0       0.78        0.65        0.71         49
      1       0.89        0.94        0.92        151

   accuracy          0.87         200
  macro avg       0.84        0.80        0.81         200
weighted avg       0.87        0.87        0.87         200

```

```

[[ 32 17]
 [ 9 142]]
F1_score = 91.61290322580645
cross_val_score= 83.9

```

In [101...]

```

# for selecting the best parameters
grid={"C":np.logspace(-4,4,20), "penalty":["l1","l2"]}
lr=LogisticRegression()
clf=GridSearchCV(lr,grid)
clf.fit(x_train,y_train)
print(clf.best_params_)

```

```
{'C': 11.288378916846883, 'penalty': 'l2'}
```

In [109...]

```

# After selecting the best parameter we need to implement them on the algorithms for check
lr=LogisticRegression(C=11.28837,penalty="l2")
lr.fit(x_train,y_train)
lr.score(x_train,y_train)
pred_lr=lr.predict(x_test)

```

```
lras=accuracy_score(y_test, pred_lr)
print('accuracy_score =', lras*100)
```

accuracy_score = 87.5

There is no much difference in the accuracy score value after performing the hyperparameter tuning by gridsearchCV

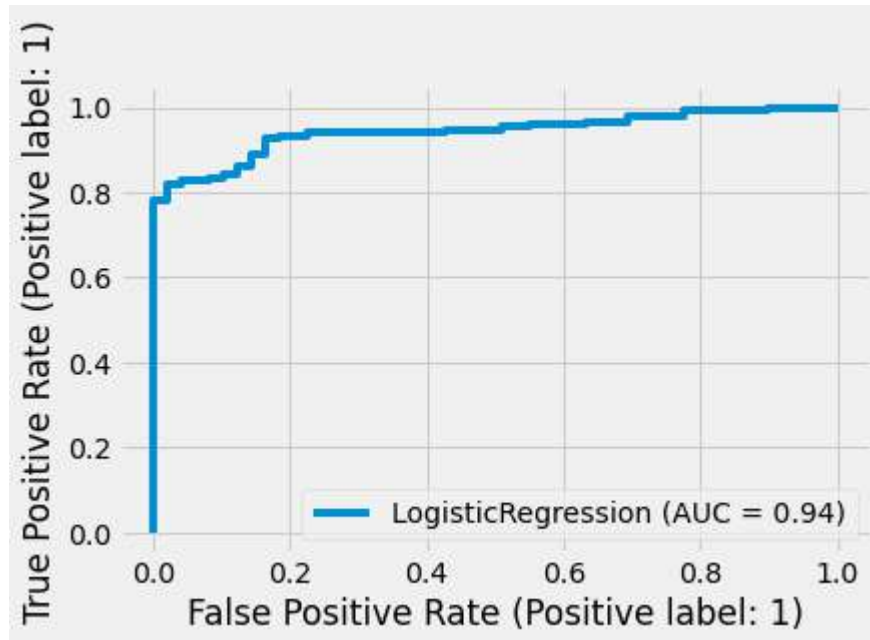
Accuracy_score is 87.5

In [126...

```
from sklearn import metrics
metrics.plot_roc_curve(lr, x_test, y_test)
metrics.roc_auc_score(y_test, pred_lr, average=None)
```

Out[126...

0.8069333693742398



By the above graph we can say that area under the curve is 94% which is very good value

CONCLUSION

- 1) area under the curve is 94percent which is very good value
- 2) The LogisticRegression giving the best accuracy value
- 3) accuracy_score, F1_score, Classification_report, Confussion_matrix and AUC value is shown in the above table
- 4) LogisticRegression is giving the best accuracy score so we need to save the LogisticRegression predicted values by using pickle